



Introduction to Data Management



Lecture #2 (Big Picture, Cont.)

Instructor: Mike Carey
mjcarey@ics.uci.edu




Announcements



- ❖ Still hanging in there? (If you do think you may drop, please decide soon so that others who are waiting can get in – *thanks...!* ☺)
- ❖ Today's plan:
 - More detail about DBMS architectures
 - Then on to logical DB design!
- ❖ Reminder:
 - Sign up on Piazza! (Only ~half have done this.)
 - HW #1 and Project Part 1 coming next week!
- ❖ Any lingering *Q's* from last time?

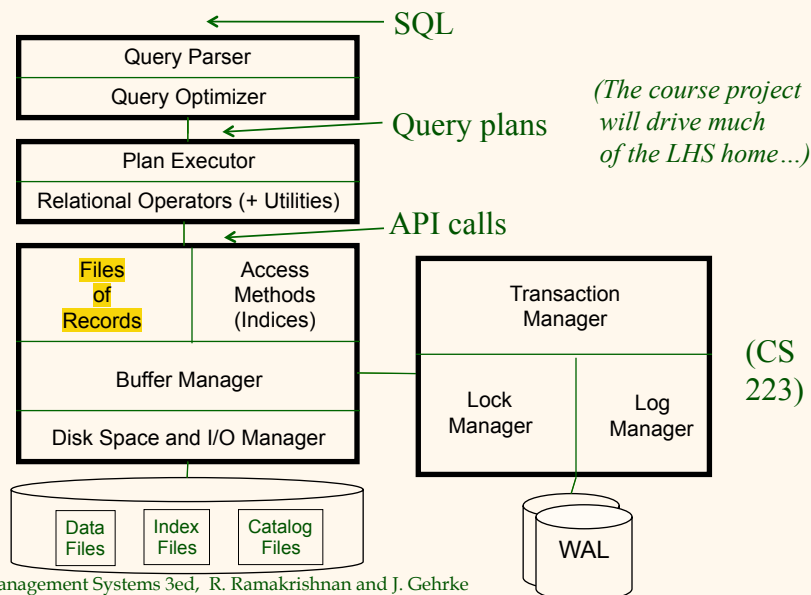
Structure of a DBMS

- ❖ A typical DBMS has a layered architecture.
- ❖ The figure does not show the concurrency control and recovery components (CS 223).
- ❖ This is one of several possible architectures; each system has its own variations.




These layers must consider concurrency control and recovery

DBMS Structure In More Detail



Components' Roles




❖ Query Parser

- Parse and analyze SQL query
- Makes sure the query is valid and talking about tables, etc., that indeed exist

❖ Query optimizer (often w/ 2 steps)


- Rewrite the query logically
 - Perform cost-based optimization
 - Goal is a “good” query plan considering
 - Physical table structures
 - Available access paths (indexes)
 - Data statistics (if known)
 - Cost model (for relational operations)
- (Cost differences can be orders of magnitude!!!)*

Components' Roles (continued)

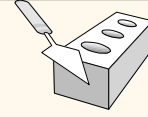


❖ Plan Executor + Relational Operators

- Runtime side of query processing
- Query plan is a tree of relational operators (drawn from the relational algebra, which you will learn about in this class)



Components' Roles (continued)




❖ Files of Records

- OSs usually have byte-stream based APIs
- DBMSs instead provide **record-based APIs**
 - Record = set of fields
 - Fields are typed
 - Records reside on pages of files

❖ Access Methods

- Index structures for access based on field values
- We'll look in a fair bit of depth at B+ tree indexes in this class (they are the most commonly used indexes across all commercial and open source systems)

Components' Roles (continued)




❖ **Buffer Manager**

- The DBMS **answer to main memory management!**
- All disk page accesses go through the buffer pool
- Buffer manager caches pages from files and indices
- "DB-oriented" page replacement scheme(s)
- Also interacts with logging (so undo/redo possible)

❖ **Disk Space** and I/O Managers

- **Manage space** on disk (pages), including extents
- Also manage I/O (sync, async, prefetch, ...)

Components' Roles (continued)




❖ System Catalog

- Info about physical data (volumes, table spaces, ...)
- Info about tables (name, columns, types, ...); also, info about their constraints, keys, etc.)
- Data statistics (e.g., value distributions, counts, ...)
- Info about indexes (types, target tables, ...)
- And so on! (Views, security, ...)

❖ Transaction Management

- ACID (Atomicity, Consistency, Isolation, Durability)
- Lock Manager for Consistency+Isolation
- Log Manager for Atomicity+Durability

Miscellany: A Few Terms



❖ Data Definition Language (DDL)

- Used to express views + logical schemas (using a syntactic form of a data model, e.g., relational)


❖ Data Manipulation Language (DML)

- Used to access and update the data in the database (again in terms of a data model, e.g., relational)

❖ Query Language (QL)

- Synonym for DML or its retrieval (i.e., data access or query) sublanguage

Miscellany (Cont'd.): Key Players



❖ Database Administrator (DBA)

- The “super user” for a database or a DBMS
- Deals with things like physical DB design, tuning, performance monitoring, backup/restore, user and group authorization management


❖ Application Developer

- Builds data-centric applications (CS122b!)
- Involved with logical DB design, queries, and DB application tools (e.g., JDBC, ORM, ...)

❖ Data Analyst or End User


- Non-expert who uses tools to interact w/ the data

A Brief History of Databases




- ❖ Pre-relational era: 1960's, early 1970's
- ❖ Codd's seminal paper: 1970
- ❖ Basic RDBMS R&D: 1970-80 (System R, Ingres)
- ❖ RDBMS improvements: 1980-85
- ❖ Relational goes mainstream: 1985-90
- ❖ Distributed DBMS research: 1980-90
- ❖ Parallel DBMS research: 1985-95
- ❖ Extensible DBMS research: 1985-95
- ❖ OLAP and warehouse research: 1990-2000
- ❖ Stream DB and XML DB research: 2000-2010
- ❖ Big data R&D: 2005-present

So Now What?




- ❖ Time to dive into the first real topic:
 - Logical DB design (ER model)
- ❖ Read the first two chapters of the book
 - Intro and ER – see the syllabus on the wiki
- ❖ Immediate to-do's for you are:
 - Get yourself signed up on Piazza
 - Stockpile sleep – no homework for you yet ☺
 - Start thinking about a project partner
- ❖ On to DB design...

Overview of Database Design



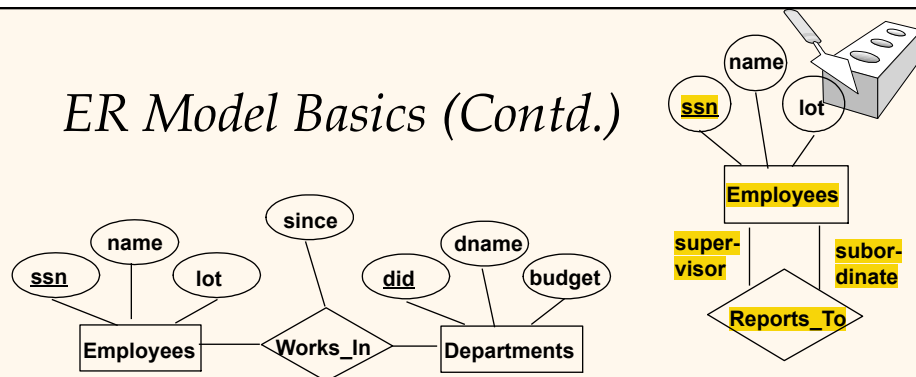
- ❖ Conceptual design: (ER Model used at this stage.)
 - What are the **entities** and **relationships** in the enterprise?
 - What information about these entities and relationships should we store in the database?
 - What are the **integrity constraints** or **business rules** that hold?
 - A database 'schema' in the ER Model can be represented pictorially (ER diagrams).
 - Can map an ER diagram into a relational schema (manually or using a design tool's automation).

ER Model Basics



- ❖ **Entity:** Real-world object distinguishable from other objects. An entity is described (in DB) using a set of **attributes**.
- ❖ **Entity Set:** A collection of similar entities.
E.g., all employees.
 - All entities in an entity set have the same set of attributes. (Until we consider ISA hierarchies, anyway... ☺)
 - Each entity set has a **key** (a unique identifier).
 - Each attribute has a **domain** (similar to a data type).

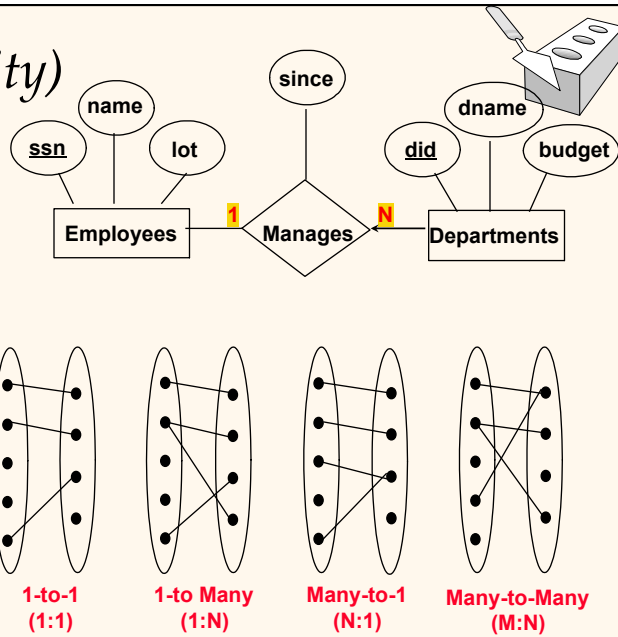
ER Model Basics (Contd.)



- ❖ **Relationship:** Association among two or more entities.
E.g., Attishoo works in Pharmacy department.
- ❖ **Relationship Set:** Collection of similar relationships.
 - An n-ary relationship set R relates n entity sets E1 ... En; each relationship in R involves entities e1 E1, ..., en En
 - Same entity set could participate in different relationship sets, or in different "roles" in same set.


Key (Cardinality) Constraints

- ❖ Consider Works_In:
An employee can work in many departments; a dept can have many employees.
- ❖ In contrast, each dept has at most one manager, according to the **key constraint** on Manages.

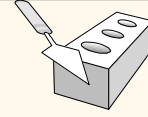


Participation Constraints

- ❖ Does every department have a manager?
 - If so, this is a **participation constraint**: the participation of Departments in Manages is said to be **total** (vs. *partial*).
 - Every Departments entity must appear in an instance of the Manages relationship
 - Ditto for both Employees and Departments for Works_In



Tune In Next Week...



❖ ... for part two of this riveting, ER modeling cliff-hanger!