



Introduction to Data Management



Lecture #3 (Conceptual DB Design)

Instructor: Mike Carey
mjcarey@ics.uci.edu




Announcements



- ❖ Today's plan:
 - More about logical DB design!
 - Basic ER concepts review and examples
 - Advanced ER concepts
- ❖ Reminders:
 - Sign up on Piazza! (Only 2/3 have done this.)
 - HW #1 and Project Part 1 coming mid-week!
 - **First quiz in discussion section tomorrow!**
- ❖ Any lingering Q's from last time?

ER Basics: Another Example



- ❖ Let's see if you can read/interpret the ER diagram above...! (☺)
 - What attributes are unique (i.e., identify their associated entity instances)?
 - What are the rules about (the much coveted) parking passes?
 - What are the rules (constraints) about professors being in departments?
 - And, what are the rules about professors heading departments?

Another Example (Cont'd.)


- ❖ Unique attributes:
 - *Professor.fac_id*, *Dept.dno*, *Parking Space.pid*
- ❖ Faculty parking:
 - 1 space/faculty, one faculty/space
 - Some faculty can bike or walk (☺)
 - Some parking spaces may be unused
- ❖ Faculty in departments:
 - Faculty may have appointments in multiple departments
 - Departments can have multiple faculty in them
 - No empty departments, and no unaffiliated faculty
- ❖ Department management:
 - One head per department (exactly)
 - Not all faculty are department heads

NOTE: These things are all “rules of the universe” that are just being modeled here!


Q: Can a faculty member head a department that he or she isn't actually in?

Weak Entities

- ❖ A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (one owner, many weak entities).
 - Weak entity set must have *total* participation in this *identifying* relationship set.
 - Dependent identifier is unique only within owner context (____), so its fully qualified key here is (ssn, dname)



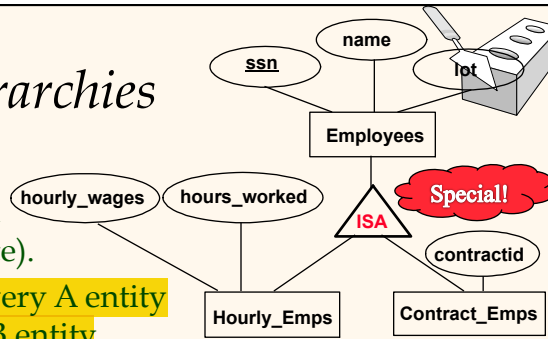
Ternary Relationships (and beyond)



- ❖ A prescription here is a 3-way relationship between a patient, a doctor, and a drug
- ❖ As modeled above, a given patient/drug combination will always be associated with one doctor
 - (General) note: Relationship key = (entity keys)

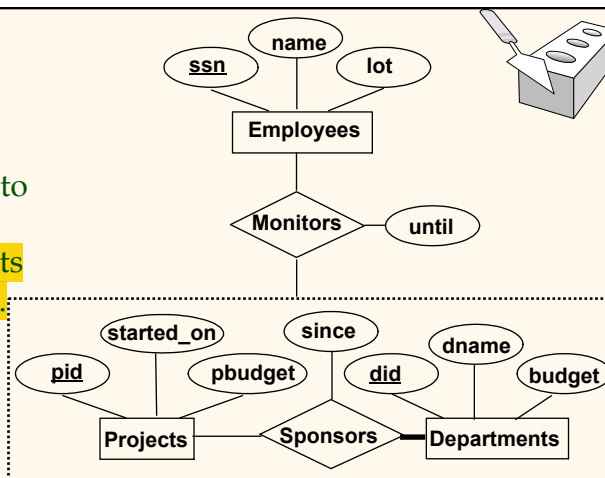
ISA ("is a") Hierarchies

- ❖ As in C++ or other PLs ER attributes are inherited (including the key attribute).
- ❖ If we declare A **ISA** B, every A entity is also considered to be a B entity.
- ❖ **Overlap constraints:** Can Joe be an Hourly_Emps as well as a Contract_Emps entity? (Allowed or **disallowed**)
 - Ex: Hourly_Emps OVERLAPS Contract_Emps (else pick 1 of the 3 types)
- ❖ **Covering constraints:** Does every Employees entity also have to be either an Hourly_Emps or a Contract_Emps entity? (Yes or **no**)
 - Ex: Hourly_Emps AND Contract_Emps COVER Employees (pick 1 of 2 vs. 1 of 3)
- ❖ Reasons for using ISA:
 - To add descriptive attributes specific to a subclass.
 - To identify subclasses that participate in a relationship.
- ❖ Design: specialization (top-down), generalization (bottom-up)




Aggregation

- ❖ Used when we have to model a relationship involving (entity sets and) a relationship set.
 - **Aggregation** allows us to treat a relationship set as an entity set for purposes of participating in (other) relationships.




- ➡ **Aggregation vs. ternary relationship:**
 - ❖ Monitors is a distinct relationship; even has its own descriptive attribute.
 - ❖ Also, can say that each sponsorship is monitored by at most one employee.


Some Advanced ER Features




- ❖ Multi-valued (vs. single-valued) attributes



- ❖ Derived (vs. stored) attributes




- ❖ Composite (vs. atomic) attributes



NOTE: Can model (two of) these using additional entity and relationship types.

Conceptual Design Using the ER Model




- ❖ Design choices:

- Should a given concept be modeled as an entity or an attribute?
- Should a given concept be modeled as an entity or a relationship?
- Characterizing relationships: Binary or ternary? Aggregation? ...

- ❖ Constraints in the ER Model:


- A lot of data semantics can (and should) be captured.
- But, not all constraints cannot be captured by ER diagrams. (Ex: Department heads from earlier...!)

Entity vs. Attribute




- ❖ Should *address* be an attribute of Employees or an **entity** (connected to Employees by a relationship)?
- ❖ Depends upon the use we want to make of address information, and the semantics of the data:
 - If we have **several** addresses per employee, *address* must be an entity (since attributes cannot be set-valued w/o advanced modeling goodies).
 - If the **structure** (city, street, etc.) is important, e.g., we want to retrieve employees in a given city, *address* must be modeled as an entity (since attribute values are atomic w/o advanced modeling goodies).
 - If the address itself is **logically separate** (e.g., the property that's located there) and refer-able, it's rightly an entity in any case!


Entity vs. Attribute (Cont'd.)



- ❖ Works_In4 does not allow an employee to work in a department for two or more periods.



- ❖ Similar to the problem of wanting to record several addresses for an employee: We want to record **several values of the descriptive attributes for each instance of this relationship**. Can be accomplished by introducing new entity set, Duration.

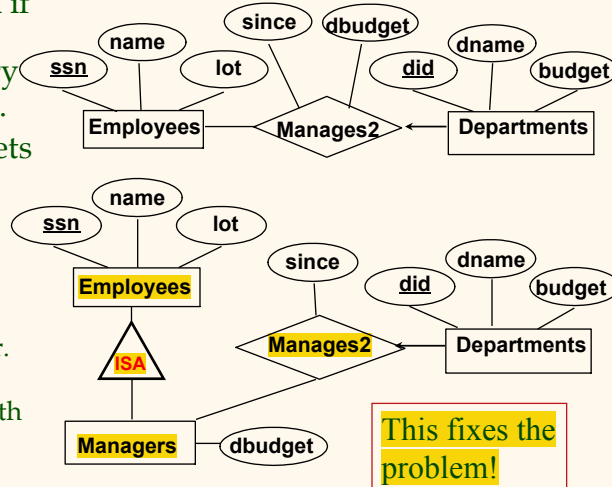


Entity vs. Relationship

- ❖ First ER diagram OK if a manager gets a separate discretionary budget for each dept.

- ❖ What if a manager gets a discretionary budget that covers all managed depts?

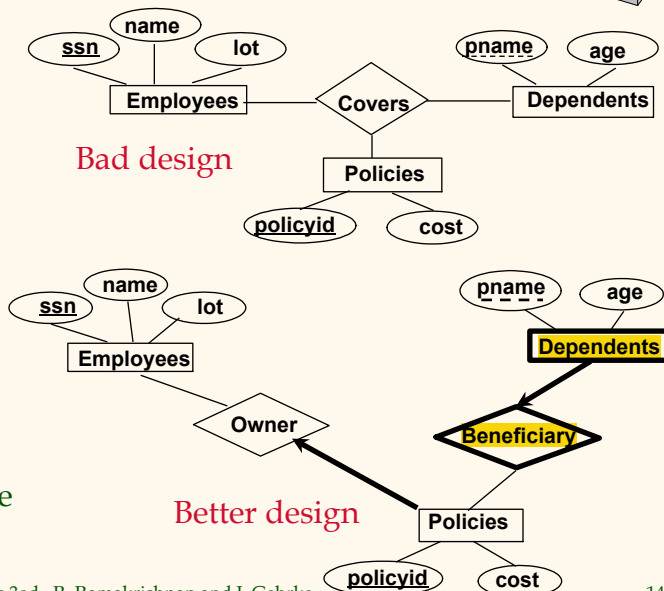
- **Redundancy:** *dbudget* stored for each dept managed by manager.
- **Misleading:** Suggests *dbudget* associated with department-mgr combination.




Binary vs. Ternary Relationships

- ❖ If each policy is owned by just **1** employee, and each **dependent** is tied to the covering policy, first diagram is inaccurate.

- ❖ **Q:** What are the additional constraints in the 2nd diagram?




Binary vs. Ternary Relationships (Cont'd.)



- ❖ Previous example illustrated a case when two binary relationships were better than one ternary relationship.
- ❖ An example in the other direction: a ternary relation **Contracts** relates entity sets **Parts**, **Departments** and **Suppliers**, and has descriptive attribute *qty*. No combination of binary relationships is an adequate substitute:
 - S “can-supply” P, D “needs” P, and D “deals-with” S does not imply that D has agreed to buy P from S.
 - How do we record *qty*?


Datbase Design Process (Flow)



- ❖ Requirements Gathering (interviews)
- ❖ Conceptual Design (using ER)
- ❖ Platform Choice (which DBMS?)
- ❖ Logical Design (for target data model)
- ❖ Physical Design (for target DBMS, workload)
- ❖ Implement (and test, of course ☺)


(Expect backtracking, iteration, and also incremental adjustments!)

Summary of Conceptual Design



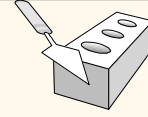
- ❖ *Conceptual design follows requirements analysis*
 - Yields a high-level description of data to be stored
- ❖ ER model popular for conceptual design
 - Constructs are expressive, close to the way people think about their applications.
- ❖ Basic constructs: **entities**, **relationships**, and **attributes** (of entities and relationships).
- ❖ Some additional constructs: **weak entities**, **ISA hierarchies**, and **aggregation**.
- ❖ Note: There are many variations on ER model (and many notations in use as well).

Summary of ER (Contd.)



- ❖ Several kinds of integrity constraints can be expressed in the ER model: *key constraints*, *participation constraints*, and *overlap/covering constraints* for ISA hierarchies. Some *foreign key constraints* are also implicit in the definition of a relationship set.
 - Some constraints (notably, *functional dependencies*) cannot be expressed in the ER model.
 - Constraints play an important role in determining the best database design for an enterprise.

Summary of ER (Contd.)



- ❖ ER design is *subjective*. There are often many ways to model a given scenario! Analyzing alternatives can be tricky, especially for a large enterprise. Common choices include:
 - Entity vs. attribute, entity vs. relationship, binary or n-ary relationship, whether or not to use ISA hierarchies, and whether or not to use aggregation.
- ❖ Ensuring good database design: resulting relational schema should be analyzed and refined further. FD information and normalization techniques are especially useful.